

**FENIL GAJJAR**

# **KUBERNETES** **DAILY TASKS**

## **KUBERNETES SERVICES**

- **COMPREHENSIVE GUIDE**
- **THEORY + PRACTICAL TASKS**
- **REAL TIME SCENARIO TASKS**
- **FOLLOW FOR MORE TASKS**

**CONTACT US**



**[fenilgajjar.devops@gmail.com](mailto:fenilgajjar.devops@gmail.com)**





**Kubernetes Services:**


**The Key to Seamless Pod**


**Communication! 🚀**



# Welcome to the Next Chapter in Our Kubernetes

## Journey – Kubernetes Services!



First of all, I want to express my **heartfelt gratitude** for the overwhelming support on my previous posts in this Kubernetes series. Your engagement, feedback, and enthusiasm make this journey even more exciting! 

So far, we've explored Kubernetes Pods – the fundamental building blocks of applications running inside Kubernetes. But have you ever wondered **how these Pods communicate with each other efficiently, even as they get replaced, rescheduled, or scaled?**  That's exactly where **Kubernetes Services** come into play!

In this post, we will **unlock the power of Kubernetes Services** – understanding how they enable seamless networking, load balancing, and service discovery across dynamic environments. Whether you're a beginner or an experienced Kubernetes practitioner, mastering Services is essential to deploying reliable and

---

scalable applications.

 **Stay tuned, keep supporting, and let's continue this learning journey together!** More **daily Kubernetes tasks** with real-world practical insights are on the way. 

 **Follow me to stay updated with hands-on Kubernetes content every day!**

Let's build, learn, and grow together! 

---

# Kubernetes Services

## What is a Kubernetes Service?

A **Service** in Kubernetes is an abstraction that provides a **stable network endpoint** to connect a set of pods. Since pods in Kubernetes are **ephemeral** (they can be created, deleted, or replaced dynamically), their IP addresses are not fixed. A Service ensures that applications can communicate with pods reliably, even if their IPs change.

### Problem Without Services:

- Each pod gets a **new IP** when restarted or rescheduled.
- Other pods **cannot reliably communicate** with a specific pod.
- There is **no built-in mechanism** for load balancing traffic across multiple pod replicas.

### Solution: Kubernetes Services

- **Provides a stable, unchanging IP & DNS name** for a group of pods.

- 
- **Automatically routes requests** to healthy pods using selectors.
  - **Load balances** traffic among multiple pod replicas.
  - **Exposes applications** within the cluster or to the external world.

## Why Do We Need Services in Kubernetes?

In Kubernetes, **Pods are ephemeral**—they can be created, deleted, or rescheduled dynamically. Since each pod gets a unique IP address that changes when it restarts, this creates challenges in communication. This is where **Kubernetes Services** come into play.

### Problem Without Services:

- **Pods have dynamic IPs** – When a pod restarts or scales, its IP changes.
- **No built-in discovery** – Other pods can't directly find or connect to a specific pod reliably.
- **Load balancing is needed** – If multiple pod replicas exist, traffic must be distributed properly.

---

## 2 How Services Solve These Problems:

A **Kubernetes Service** provides a **stable IP and DNS name** for a group of pods, making them accessible reliably.

### ✓ Key Benefits of Services:

1. **Stable Networking** – Services ensure a consistent way to communicate with pods, even if pod IPs change.
2. **Automatic Load Balancing** – Distributes traffic across multiple pod replicas.
3. **Service Discovery** – Kubernetes assigns a DNS name (`my-service.default.svc.cluster.local`) so other pods can find it easily.
4. **External Access** – Services allow external users to access applications running inside the cluster.

## 3 Types of Kubernetes Services

**Service Type**

**Use Case**

---

<b>ClusterIP</b> <b>(Default)</b>	Exposes the service <b>internally</b> within the cluster. Used for pod-to-pod communication.
<b>NodePort</b>	Exposes the service on a <b>static port on each node</b> . Used for basic external access.
<b>LoadBalancer</b>	Integrates with a cloud provider's <b>load balancer</b> (AWS, GCP, Azure) for external access.
<b>ExternalName</b>	Maps a service to an <b>external domain name</b> (e.g., <code>example.com</code> ).

## 4 . How Services Work with DNS in Kubernetes

Kubernetes assigns a **DNS name to each service**, making communication simple.

For example:

- A pod can access the service using `http://web-service` instead of using pod IPs.
- The actual DNS name is `web-service.default.svc.cluster.local`, making it discoverable inside the cluster.



---


# ClusterIP Service in Kubernetes

## 1 What is a ClusterIP Service?

A **ClusterIP** service in Kubernetes is the **default type of service** that enables internal communication between different pods within a Kubernetes cluster. It assigns a **stable internal IP** that is accessible only from within the cluster but **not from the outside world**.

### Why Do We Need ClusterIP?

- **Pods are ephemeral** – Their IP addresses change when restarted.
- **Pods need to communicate** with each other inside the cluster.
- **Stable Networking** – ClusterIP ensures a consistent way to access pods.

 **Use Case:** Internal microservices communication, backend-to-database connections, etc.

## 2 How Does a ClusterIP Service Work?

1. A **Service** is created with type **ClusterIP** and a **label selector** to target specific pods.

- 
2. Kubernetes assigns a **stable virtual IP (ClusterIP)** to the service.
  3. Any pod inside the cluster can access the service using:
    - **ClusterIP**
    - **Service name (via DNS)**
  4. Kubernetes **automatically load balances** traffic across the selected pods.

 **Important:** The ClusterIP is only accessible **inside the cluster**.

### Example: Creating a ClusterIP Service

#### Step 1: Create a Deployment

Let's deploy an **Nginx-based web app** running in multiple pods.

**apiVersion: apps/v1**

**kind: Deployment**

**metadata:**

**name: web-app**

**spec:**

**replicas: 3**

---

**selector:**

**matchLabels:**

**app: web-app**

**template:**

**metadata:**

**labels:**

**app: web-app**

**spec:**

**containers:**

**- name: nginx**

**image: nginx**

**ports:**

**- containerPort: 80**

---

## Step 2: Create a ClusterIP Service

Now, let's expose this deployment using a **ClusterIP service**.

**apiVersion: v1**

**kind: Service**

**metadata:**

**name: web-service**

**spec:**

**selector:**

**app: web-app # Targets pods with this label**

**ports:**

**- protocol: TCP**

**port: 80 # Port exposed by the service**

**targetPort: 80 # Port inside the pod**

**type: ClusterIP**

---

## 4 Accessing a ClusterIP Service

Once the service is created, Kubernetes assigns it a **ClusterIP**, which can be seen using:

```
kubectl get services
```

 **Example Output:**

NAME	TYPE	CLUSTER-IP	PORT(S)	AGE
web-service	ClusterIP	10.96.15.200	80/TCP	5m

**Ways to Access the Service:**

 **Using ClusterIP:**

```
curl http://10.96.15.200:80
```

 **Using Kubernetes DNS Name (Preferred):**

```
curl http://web-service
```

- 
- DNS Name: `web-service.default.svc.cluster.local`
  - Automatically resolves to `10.96.15.200` within the cluster.

## 5 How Kubernetes Handles ClusterIP Internally

Kubernetes uses **iptables** or **IPVS** to route traffic to the backend pods.

### ♦ Steps:

1. A request to `http://web-service` is received.
2. The request is routed via `kube-proxy`.
3. It forwards traffic to one of the **healthy pods** running `web-app`.
4. **Load balancing** happens automatically.

### 📌 Round-Robin Load Balancing:

If `web-service` has 3 pods running, requests will be distributed evenly among them.

---

## 6 Real-World Use Cases of ClusterIP

### ✓ Microservices Communication

- A frontend service (`frontend-service`) calls a backend API (`backend-service`).
- Backend interacts with a database service (`db-service`).

### ✓ Database Access Inside Cluster

- A MySQL pod is exposed as `mysql-service`.
- Backend pods use `mysql-service:3306` to connect instead of direct pod IPs.

### ✓ Internal Caching Systems

- Redis or Memcached can be exposed via ClusterIP (`redis-service`).
- Other applications access it for caching.

---

## 7 ClusterIP vs. Other Service Types

Service Type	Accessibility	Use Case
<b>ClusterIP</b> (default)	Only inside the cluster	Internal communication (microservices, databases)
<b>NodePort</b>	Exposes on a fixed port on each node	Direct external access (debugging, testing)
<b>LoadBalancer</b>	Exposes via cloud load balancer	Production applications requiring public access
<b>ExternalName</b>	Maps to an external DNS	Connecting to external databases (AWS RDS, etc.)

 Without ClusterIP, internal communication in Kubernetes would be unreliable! 



---

# Real-Life Scenario: ClusterIP Service in Kubernetes

## Scenario: E-commerce Application with Microservices

Imagine you are deploying an **e-commerce platform** in Kubernetes. The architecture consists of multiple microservices:

1. **Frontend Service** (React/Angular)
2. **Backend API Service** (Node.js/Java)
3. **Database Service** (MySQL)
4. **Redis Cache** (For faster lookups)

Since the backend, database, and caching system **only need to communicate internally**, we use **ClusterIP services** to expose them.

## Step 1: Deploy Backend API Service

The backend API service is responsible for handling requests from the frontend and communicating with the database.

---

## Backend Deployment

```
apiVersion: apps/v1

kind: Deployment

metadata:
  name: backend-api

spec:
  replicas: 3

  selector:
    matchLabels:
      app: backend-api

  template:
    metadata:
      labels:
        app: backend-api

    spec:
```

---

```
containers:
```

```
- name: backend
```

```
  image: mycompany/backend-service:latest
```

```
  ports:
```

```
    - containerPort: 8080
```

## Backend ClusterIP Service

This exposes the backend service **only inside the cluster**.

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: backend-service
```

```
spec:
```

```
  selector:
```

```
    app: backend-api
```

```
  ports:
```

---

- protocol: TCP

port: 80

targetPort: 8080

type: ClusterIP

### How It Works?

- Kubernetes assigns a **ClusterIP** (e.g., `10.96.15.100`).
- Any pod inside the cluster can access it using

`http://backend-service`.

## Step 2: Deploy MySQL Database

Since the database should **only be accessible inside the cluster**, we use another **ClusterIP service**.

### MySQL Deployment

apiVersion: apps/v1

kind: Deployment

---

metadata:

name: mysql-db

spec:

replicas: 1

selector:

matchLabels:

app: mysql-db

template:

metadata:

labels:

app: mysql-db

spec:

containers:

- name: mysql

image: mysql:5.7

env:

---

```
- name: MYSQL_ROOT_PASSWORD
```

```
  value: "mypassword"
```

```
- name: MYSQL_DATABASE
```

```
  value: "ecommerce"
```

```
ports:
```

```
- containerPort: 3306
```

## MySQL ClusterIP Service

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: mysql-service
```

```
spec:
```

```
  selector:
```

```
    app: mysql-db
```

---

ports:

- protocol: TCP

port: 3306

targetPort: 3306

type: ClusterIP

### How It Works?

- Backend API accesses MySQL using `mysql-service:3306` instead of a changing pod IP.
- The database is **protected from external access**.

## Step 3: Deploy Redis Cache

The backend API needs **fast access** to Redis, so we create another **ClusterIP service**.

### Redis Deployment

`apiVersion: apps/v1`

---

kind: Deployment

metadata:

name: redis-cache

spec:

replicas: 1

selector:

matchLabels:

app: redis-cache

template:

metadata:

labels:

app: redis-cache

spec:

containers:

- name: redis

image: redis:latest



---

ports:

- containerPort: 6379

### Redis ClusterIP Service

apiVersion: v1

kind: Service

metadata:

- name: redis-service

spec:

- selector:

- app: redis-cache

- ports:

- protocol: TCP

- port: 6379

- targetPort: 6379

- type: ClusterIP

---

## How It Works?

- Backend API connects to `redis-service:6379`.
- The Redis service is **not accessible from outside the cluster**.

## Step 4: Frontend Access

Since the **frontend needs to be accessible externally**, we use a **LoadBalancer or NodePort service** for it.

### Frontend Service (LoadBalancer)

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: frontend-service
```

```
spec:
```

```
  selector:
```

---

app: frontend

ports:

- protocol: TCP

port: 80

targetPort: 3000

type: LoadBalancer

 **Frontend accesses services as:**

- backend-service:80
- mysql-service:3306
- redis-service:6379

---

# NodePort Service in Kubernetes

## What is a NodePort Service?

A **NodePort service** in Kubernetes exposes an application **on every node in the cluster** using a specific port (called a **NodePort**). This allows external traffic to access the service **directly via any node's IP and the assigned port**.

Unlike a **ClusterIP service**, which is only accessible within the cluster, a **NodePort service** makes the application accessible from outside the cluster.

## How NodePort Works?

1. **Kubernetes assigns a static port (30000-32767)** to the service.
2. **Each node in the cluster listens on this port** and forwards traffic to the service.

**Requests can be made using:**

`http://<NodeIP>:<NodePort>`

---

3. where:

- `<NodeIP>` is the IP of any node in the cluster.
- `<NodePort>` is the assigned port (e.g., `30007`).

## Example: Exposing a Web Application using NodePort

Let's say we have a **Node.js web application** running inside a Kubernetes pod. We want to expose it externally using **NodePort**.

### Step 1: Create a Deployment

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: web-app
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

---

```
    app: web-app

template:

  metadata:

    labels:

      app: web-app

  spec:

    containers:

      - name: web-app

        image: mycompany/web-app:latest

        ports:

          - containerPort: 8080
```

- This will create **two pods** running the web application on port **8080**.

---

## Step 2: Expose the Application Using NodePort

```
apiVersion: v1

kind: Service

metadata:
  name: web-app-service

spec:
  type: NodePort

  selector:
    app: web-app

  ports:
    - protocol: TCP

      port: 80          # ClusterIP Port

      targetPort: 8080 # The port in the container

      nodePort: 30007  # Exposed on each node (between
30000-32767)
```

---

### Step 3: Access the Application

After applying the YAML files, **Kubernetes will expose the service on every node's IP at port 30007.**

You can access the web application using:

`http://<NodeIP>:30007`

- If you have multiple nodes, you can use **any** node's IP.
- Traffic coming to **30007** is forwarded to the **ClusterIP service**, which routes it to the **backend pods**.

### Use Cases of NodePort

- **Exposing applications during development or testing.**
- **Accessing applications from outside the cluster** without using an Ingress.
- **Running internal services that need to be accessed without cloud load balancers.**



---

## Limitations of NodePort



**Not ideal for production** because:

- The port range **(30000-32767)** is limited.
- Requires knowing **node IPs manually**.
- Not suitable for large-scale applications with multiple services.

For production, **LoadBalancer or Ingress** is preferred.

---

# Real-Life Scenario of NodePort Service

## Scenario: Exposing an Internal Web Dashboard for a Private Network

Imagine a company has a **Kubernetes cluster running on-premises** (not in the cloud) and hosts an **internal web dashboard** for monitoring system performance. The dashboard is only meant to be accessed by employees within the company's private network.

Since this is an **on-premises setup**, there is **no cloud provider** to provide a **LoadBalancer** service. Instead, the company decides to use **NodePort** to expose the dashboard to employees.

## How NodePort Helps in This Scenario?

1. The dashboard application is running inside Kubernetes pods.
2. A **NodePort service** exposes the dashboard on port **32000**.

Employees within the **private network** can access it using:

**`http://<AnyNodeIP>:32000`**

- 
3. The NodePort service forwards the request to the **correct pod** running the dashboard.

## Example NodePort YAML for Internal Web Dashboard

### Step 1: Create the Deployment

```
apiVersion: apps/v1

kind: Deployment

metadata:

  name: dashboard

spec:

  replicas: 2

  selector:

    matchLabels:

      app: dashboard

  template:

    metadata:
```

---

```
  labels:

    app: dashboard

spec:

  containers:

    - name: dashboard

      image: mycompany/dashboard:latest

      ports:

        - containerPort: 8080
```

- This deployment creates **two replicas** of the web dashboard application.
- The application runs on **port 8080** inside the container.

## Step 2: Expose the Dashboard Using NodePort

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

---

```
name: dashboard-service
```

```
spec:
```

```
type: NodePort
```

```
selector:
```

```
  app: dashboard
```

```
ports:
```

```
  - protocol: TCP
```

```
    port: 80          # ClusterIP port
```

```
    targetPort: 8080 # Pod container port
```

```
    nodePort: 32000  # Exposed on all nodes
```

- The service listens on **port 32000** on **every node**.
- Requests received on **32000** are forwarded to **80**, which then routes to the **pods running the dashboard on port 8080**.

### Step 3: Access the Dashboard

---

Employees can access the dashboard from their browsers by visiting:

`http://<NodeIP>:32000`

- `<NodeIP>` can be **any Kubernetes worker node IP** in the cluster.
- The request reaches `32000`, which **forwards traffic to the pods** running the dashboard.

## Why Use NodePort in This Case?

✓ **No need for a cloud-based LoadBalancer** (as this is an on-premises cluster).

✓ **Internal employees can access it via private network without extra networking setup.**

✓ **Simple way to expose services within an organization.**

However, for **external access over the internet**, a **LoadBalancer or Ingress** is a better choice. 🚀

---

## Default Kubernetes Service in the Cluster:

### kubernetes Service

When a **Kubernetes cluster** is created, a default service named **kubernetes** is automatically created.

#### 1 What is the **kubernetes** Service?

- It is a **ClusterIP service** created by Kubernetes itself.
- It provides a **stable entry point** to the Kubernetes API server.
- It allows **internal cluster components** (like Pods) to communicate with the **Kubernetes API** without knowing its actual IP.

#### 2 Why Does Kubernetes Create This Service by Default?

The **kubernetes** service exists to: ☒ Allow Pods to **discover and interact with the Kubernetes API**.

☒ Enable internal components like **Controllers, Schedulers, and Custom Operators** to interact with the API server.

☒ Provide a **fixed internal DNS name**

---

(**kubernetes.default.svc.cluster.local**) that always resolves to the API server.

### 3 Details of the Default **kubernetes** Service

#### ● Checking the Service

You can verify the default service using:

```
kubectl get svc
```

This will show:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP

10d

- The **kubernetes** service is of type **ClusterIP**.
- It has a fixed internal IP (e.g., **10.96.0.1**).
- It listens on port **443 (HTTPS)** for API requests.



---

#### 4 How Pods Use This Service?

Every pod in the cluster can access the API using:

```
curl -k https://kubernetes.default.svc.cluster.local
```

OR

```
curl -k https://10.96.0.1
```

#### Example Use Case:

Suppose a pod wants to list all services in the cluster. It can send an API request like this:

```
curl -k -H "Authorization: Bearer $(cat  
/var/run/secrets/kubernetes.io/serviceaccount/token)" \  
  
https://kubernetes.default.svc/api/v1/services
```

- This request reaches the Kubernetes API server via the **default service**.
- The API server processes the request and returns the list of services.

---

## 5 How is This Service Created?

- Kubernetes **automatically creates** this service during cluster setup.
- It uses a **virtual IP (ClusterIP)** to route traffic to the API server.
- It is managed by **kube-apiserver**.

## 6 Can You Delete the **kubernetes** Service?



**No, you should not delete it!**

- It is **critical** for internal Kubernetes operations.
- Many components (like controllers and operators) rely on it.
- Without it, **internal pods cannot communicate with the API server**.

However, if accidentally deleted, you can **recreate it manually**:

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: kubernetes
```

---

```
namespace: default

spec:

  type: ClusterIP

  clusterIP: 10.96.0.1 # Ensure this matches your cluster's
settings

  ports:

    - port: 443

      targetPort: 443

  selector:

    component: apiserver
```

Then, apply it using:

```
kubectl apply -f kubernetes-service.yaml
```

---

## Default ClusterIP Kubernetes Service: Why Is It Useful?

In every Kubernetes cluster, a **default ClusterIP service** is automatically created to enable essential **internal communication** between cluster components. One such default service is **kubernetes**, which provides an internal DNS entry (**kubernetes.default.svc.cluster.local**) that allows pods to communicate with the Kubernetes API server.

### ◆ Why Is the Default ClusterIP Service Useful?

#### 1 Kubernetes API Access for Internal Components

📌 **Scenario:** Pods running in the cluster need to interact with the Kubernetes API for operations like scaling, retrieving configurations, or updating resource status.

✅ **How it helps:** The **kubernetes** service provides a stable internal IP (**ClusterIP**), so workloads can securely communicate with the API server without needing to hardcode its IP address.

#### 💡 Example:

A pod using a service account to fetch information about available nodes:

---

```
curl -k -H "Authorization: Bearer $(cat  
/var/run/secrets/kubernetes.io/serviceaccount/token)" \  
https://kubernetes.default.svc/api/v1/nodes
```

♦ **Why?** This allows applications to discover and interact with cluster resources dynamically.

## 2 Service Discovery for Cluster Components

📌 **Scenario:** Internal services (like CoreDNS, monitoring tools, or admission controllers) need to find and communicate with the Kubernetes API.

✅ **How it helps:** The `kubernetes` service ensures that no matter where a pod is scheduled, it can always find the API server at `kubernetes.default.svc`.

### 💡 Example:

- CoreDNS uses the Kubernetes API to resolve service names.
- Metrics servers query the API to collect resource usage data.

## 3 Secure Cluster-Level Communication

📌 **Scenario:** Applications or system services need **authenticated** and **secure**

---

access to cluster data.

✅ **How it helps:** The service provides **TLS-secured access** to the API, enforcing RBAC (Role-Based Access Control) policies.

💡 **Example:**

- A monitoring agent in a pod queries Kubernetes for node health.
- A CI/CD pipeline interacts with Kubernetes to deploy applications.

#### 4 Custom Controllers & Operators Need API Access

📌 **Scenario:** Custom controllers or operators running in the cluster need to **watch Kubernetes resources** and react to changes.

✅ **How it helps:** The service allows controllers to interact with the API server via the **watch** mechanism.

💡 **Example:**

A **custom Kubernetes operator** that automatically scales resources based on a database load.

```
apiVersion: batch/v1
```

```
kind: Job
```

---

metadata:

name: scale-job

spec:

template:

spec:

containers:

- name: scaler

image: myorg/scaler:latest

command: ["python", "/scripts/scale\_resources.py"]

env:

- name: KUBERNETES\_SERVICE\_HOST

value: "kubernetes.default.svc"

- ♦ **Why?** The operator doesn't need to worry about API server IP changes.

---

## ◆ When Does This Service Become Critical?

- ✓ When internal Kubernetes components (like CoreDNS, kube-proxy, or controllers) need to communicate with the API server.
- ✓ When applications inside the cluster need to interact with Kubernetes dynamically.
- ✓ When running CI/CD pipelines that deploy applications using `kubectl` or Kubernetes API calls.
- ✓ When monitoring or logging tools like Prometheus and Fluentd need to query Kubernetes metadata.



---

## LoadBalancer Service in Kubernetes (Using Kops)

### ◆ What Is a LoadBalancer Service?

A **LoadBalancer** service in Kubernetes automatically provisions an **external load balancer** to expose your application to the internet or external networks. It integrates with cloud provider load balancers (e.g., AWS ELB, GCP Load Balancer, Azure Load Balancer).

If you are using **Kops** to create a Kubernetes cluster, it means you're likely running on **AWS**, and Kubernetes will provision an **AWS Elastic Load Balancer (ELB)** when you create a **LoadBalancer** service.

### ◆ How LoadBalancer Works in a Kops Cluster (AWS

### Example)

#### 1 You Create a LoadBalancer Service

- When you define a **LoadBalancer** service in Kubernetes, Kops automatically provisions an **AWS ELB**.
- The ELB forwards traffic to worker nodes where the app is running.

---

## 2 Traffic Is Distributed to Worker Nodes

- The ELB sends traffic to worker nodes on the specified ports.
- Kubernetes **kube-proxy** ensures that the traffic is forwarded to the correct pod.

## 3 Pods Receive Requests via NodePort Service

- Internally, Kubernetes **creates a NodePort service** and assigns a port on each worker node.
- The ELB routes external traffic to this **NodePort**, which then directs traffic to the correct pods.

### ◆ Example: LoadBalancer Service in Kops (AWS ELB

#### Example)

Here's how you define a **LoadBalancer** service for an Nginx deployment:

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

---

```
name: nginx-service
```

```
annotations:
```

```
    service.beta.kubernetes.io/aws-load-balancer-type: "nlb"
```

```
# Use Network Load Balancer (NLB)
```

```
spec:
```

```
  type: LoadBalancer
```

```
  selector:
```

```
    app: nginx
```

```
  ports:
```

```
    - protocol: TCP
```

```
      port: 80           # External Port
```

```
      targetPort: 8080  # Container Port
```

---

## ◆ What Happens When You Apply This YAML?

### 1 Kubernetes Requests an AWS ELB

- Kops provisions an ELB in AWS.
- The ELB gets a **public IP / DNS name**, making the service accessible externally.

### 2 Traffic Flows Through the Load Balancer

- The ELB forwards traffic to the worker nodes on a dynamically assigned **NodePort**.
- Kubernetes then routes traffic to the correct pods using **kube-proxy**.

### 3 DNS Name for Access

AWS assigns a DNS name like:

**a1b2c3d4e5-1234567890.us-east-1.elb.amazonaws.com**

- You can access the service using this DNS name instead of manually configuring an ingress.

---

## ◆ Key Features & Benefits

✓ **Automatic External Access** – No need for manual configuration; Kubernetes provisions an ELB automatically.

✓ **High Availability** – ELB distributes traffic across multiple nodes, ensuring redundancy.

✓ **Integration with Cloud Providers** – Works seamlessly with AWS, GCP, Azure, etc.

✓ **Supports Layer 4 (NLB) & Layer 7 (ALB) Load Balancers** – You can choose between TCP (NLB) or HTTP/HTTPS (ALB) balancing.

## ◆ Use Cases for LoadBalancer Services

✓ **Publicly Exposing Web Applications** – Websites, APIs, or services that need external access.

✓ **Multi-Region or Multi-Zone Deployments** – When you want a **highly available** setup.

✓ **Integrating with DNS & CDN** – Can be used with Route 53 and CloudFront for global load balancing.

---

✓ **Secure HTTPS Access** – Supports TLS termination when using an Application Load Balancer (ALB).

## Troubleshooting & Best Practices

- ◆ **Check the Created Load Balancer:**

```
kubectl get svc nginx-service
```

- ◆ **Find External IP/DNS of Load Balancer:**

```
kubectl describe svc nginx-service
```

- ◆ **Use Ingress for Advanced Routing** – Instead of directly using a LoadBalancer, consider **Ingress controllers** for better path-based routing and TLS termination.

- ◆ **Enable Connection Draining** – Ensures smooth traffic shifting when pods are replaced or scaled.

---

## ◆ Understanding the YAML Configuration

This YAML defines:

- 1 A Pod (**nginx**) running the **nginx** container.
- 2 A Service (**nginx-service**) exposing the **nginx** Pod.

### ◆ Key Points in the Pod Definition

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

```
  labels:
```

```
    app.kubernetes.io/name: proxy # Label used for service  
selection
```

```
spec:
```

---

```
containers:
```

```
- name: nginx
```

```
  image: nginx:stable
```

```
  ports:
```

```
    - containerPort: 80
```

```
      name: http-web-svc # Named port inside the
```

```
container
```

- The **Pod is labeled** as `app.kubernetes.io/name: proxy`, which the Service uses to **select and route traffic**.
- The **container exposes port 80**, but importantly, the port is given a **name** (`http-web-svc`).
- Named ports allow **flexibility** when referring to ports in Services.

#### ◆ Key Points in the Service Definition

```
apiVersion: v1
```

```
kind: Service
```



---

metadata:

name: nginx-service

spec:

selector:

app.kubernetes.io/name: proxy # Selects Pods with this  
label

ports:

- name: name-of-service-port # Name for the Service port  
(optional)

protocol: TCP

port: 80 # Exposed port of the Service

targetPort: http-web-svc # Refers to the named  
containerPort inside Pods

- The Service **selects Pods** with the label `app.kubernetes.io/name: proxy`.

- 
- The Service **exposes port 80 externally** and forwards traffic to the **named port (http-web-svc)** in the selected Pods.
  - Using named ports (**targetPort: http-web-svc**) instead of hardcoded numbers provides **flexibility when changing ports** in the future.

## ◆ Why Is This Flexible?

### Scenario: Changing the Container's Port Without Breaking the Service

Imagine you **update the backend software** in the `nginx` container, and now it **listens on port 8080 instead of 80**.

Without named ports, you'd need to manually update the **targetPort** in the Service. However, with named ports, you **only update the Pod's containerPort**, and the Service automatically routes traffic correctly.

#### Updated Pod with New Port

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

---

labels:

app.kubernetes.io/name: proxy

spec:

containers:

- name: nginx

image: nginx:stable

ports:

- containerPort: 8080 # Changed from 80 to 8080

name: http-web-svc # Still the same name

- **No need to change the Service definition** because it still refers to `targetPort: http-web-svc`.
- Kubernetes automatically updates the routing based on the new Pod configuration.

---

## ◆ Real-World Benefits of This Approach

✓ **Backward Compatibility:** Clients using the Service don't break when Pod ports change.

✓ **Seamless Rolling Updates:** You can roll out new versions of an app with different ports without breaking existing traffic.

✓ **Multi-Port Flexibility:** If different backend Pods expose different ports, a Service can still route to them using named ports.

---

## Example: Multiple Pods with Different Ports in the Same Service

### ◆ Scenario

Imagine a **microservices-based backend** where different versions of an application (e.g., v1 and v2) **run in separate Pods**, and each version **exposes a different port**. However, we want a **single Service** to route traffic to all of them **without breaking clients**.

### ◆ Solution Using Named Ports

Instead of **hardcoding port numbers**, we use **named ports** in the Pods and reference them in the Service.

### Step 1: Define Two Pods (v1 & v2) with Different Ports

```
apiVersion: v1
```

```
kind: Pod
```

---

metadata:

name: backend-v1

labels:

app.kubernetes.io/name: backend

version: v1

spec:

containers:

- name: backend

image: my-backend:v1

ports:

- containerPort: 8080 # v1 runs on 8080

name: backend-port

---

apiVersion: v1

kind: Pod

---

metadata:

name: backend-v2

labels:

app.kubernetes.io/name: backend

version: v2

spec:

containers:

- name: backend

image: my-backend:v2

ports:

- containerPort: 9090 # v2 runs on 9090

name: backend-port

- **backend-v1 Pod** listens on **port 8080**.
- **backend-v2 Pod** listens on **port 9090**.
- **Both Pods use the same named port (backend-port)** for consistency.

---

## Step 2: Define a Single Service to Route Traffic

```
apiVersion: v1

kind: Service

metadata:
  name: backend-service

spec:
  selector:
    app.kubernetes.io/name: backend # Matches both v1 & v2
  ports:
    - name: service-port
      protocol: TCP
      port: 80 # Service exposes port 80
      targetPort: backend-port # Routes to named port in Pods
```



- 
- The Service **routes requests coming to port 80** to the Pods' **backend-port**.
  - Since each Pod has **different containerPorts (8080, 9090)** but **the same named port (backend-port)**, the Service **automatically maps them correctly**.
  - Clients can access the backend **without knowing the exact backend Pod ports**.

#### ♦ What Happens Internally?

- 1 A user makes a request to **backend-service** on **port 80**.
- 2 The Service **forwards the request to any running Pod (backend-v1 or backend-v2)**.
- 3 Kubernetes **automatically resolves the correct port (8080 for v1, 9090 for v2)** using the named port **backend-port**.

---

## Understanding the Multi-Port Service in Kubernetes

This Kubernetes **Service** definition exposes **two ports (80 and 443)** and forwards traffic to **two different target ports (9376 and 9377)** on selected Pods.

### ◆ Breakdown of the YAML

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: my-service
```

```
spec:
```

```
  selector:
```

```
    app.kubernetes.io/name: MyApp # Selects Pods with this  
label
```

```
  ports:
```

```
    - name: http
```

---

```
protocol: TCP

port: 80          # Service listens on port 80

targetPort: 9376 # Forwards traffic to 9376 on
selected Pods
```

```
- name: https

protocol: TCP

port: 443          # Service listens on port 443

targetPort: 9377 # Forwards traffic to 9377 on
selected Pods
```

- **Selector:** Matches all Pods labeled as **app.kubernetes.io/name: MyApp**.
- **Ports:**
  - **Port 80** (HTTP) → forwards requests to **9376** inside Pods.
  - **Port 443** (HTTPS) → forwards requests to **9377** inside Pods.
- **Traffic Distribution:** Any request coming to **port 80 or 443 of the Service** gets routed to a **Pod running MyApp**.

---

## ♦ Example Use Case: Exposing a Web Application

Imagine we have a **web application with a backend API** running inside Kubernetes.

### Step 1: Create Two Pods (with Different Ports)

```
apiVersion: v1

kind: Pod

metadata:

  name: myapp-pod

  labels:

    app.kubernetes.io/name: MyApp

spec:

  containers:

    - name: web-server

      image: nginx

      ports:
```

---

```
- containerPort: 9376 # Handles HTTP requests

- name: api-server

  image: my-api

  ports:

    - containerPort: 9377 # Handles HTTPS requests
```

- The **web-server container** (e.g., `nginx`) listens on **port 9376**.
- The **API server container** listens on **port 9377** for secure traffic.

## Step 2: Create the Multi-Port Service

```
apiVersion: v1

kind: Service

metadata:

  name: myapp-service

spec:

  selector:
```

---

```
app.kubernetes.io/name: MyApp
```

```
ports:
```

```
- name: http
```

```
  protocol: TCP
```

```
  port: 80
```

```
  targetPort: 9376 # Web server
```

```
- name: https
```

```
  protocol: TCP
```

```
  port: 443
```

```
  targetPort: 9377 # API server
```

- When users access **http://myapp-service:80**, requests go to the **web server** on port **9376**.
- When users access **https://myapp-service:443**, requests go to the **API server** on port **9377**.

---

### ♦ Why Is This Useful?

✓ **Single Service, Multiple Ports:** We don't need separate Services for HTTP and HTTPS.

✓ **Load Balancing:** Kubernetes distributes traffic between multiple Pods running the same app.

✓ **Simplified Network Configuration:** Clients only need to know the Service name (`myapp-service`), not the backend port details.

---

# Headless Services in Kubernetes

A **Headless Service** in Kubernetes is a type of Service that **does not provide load balancing or a single cluster IP**, but instead directly exposes the **individual Pods' IP addresses**. This is useful when applications need direct access to each Pod, such as databases or StatefulSets.

## ◆ How Does a Headless Service Work?

- A normal Kubernetes **Service (ClusterIP)** assigns a single IP address and load balances traffic to different Pods.
- A **Headless Service (ClusterIP: None)** does **not** assign a single IP.
- Instead, it **returns individual Pod IPs** when queried, allowing clients to connect directly to specific Pods.

## ◆ Why Use a Headless Service?

✓ **Direct Communication** - Useful for databases like **MySQL, MongoDB, and Cassandra**, where clients need to communicate with specific Pods.

✓ **Stateful Applications** - When a Pod has persistent data and cannot be



---

randomly load-balanced.

✓ **Service Discovery** - Applications can discover and interact with each Pod individually.

## ◆ Example: Headless Service for a Database

Imagine a **MongoDB StatefulSet** where each Pod needs to be directly accessible.

### 1 Create a Headless Service

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: mongo-service
```

```
spec:
```

```
  clusterIP: None # Makes it headless
```

```
  selector:
```

```
    app: mongo
```

```
  ports:
```

---

- port: 27017

targetPort: 27017

♦ **clusterIP: None** means no virtual IP is assigned, and DNS will resolve to individual Pod IPs.

## 2 Create a StatefulSet with MongoDB

apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mongo

spec:

serviceName: mongo-service # Links to the Headless  
Service

replicas: 3 # Three MongoDB replicas

selector:

---

```
matchLabels:

  app: mongo

template:

  metadata:

    labels:

      app: mongo

  spec:

    containers:

      - name: mongo

        image: mongo

        ports:

          - containerPort: 27017
```

♦ The **StatefulSet** ensures that **Pods get stable hostnames** like:

- `mongo-0.mongo-service.default.svc.cluster.local`
- `mongo-1.mongo-service.default.svc.cluster.local`

- 
- `mongo-2.mongo-service.default.svc.cluster.local`

## ◆ How Does It Work?

- Instead of a single **mongo-service** IP, DNS resolves to **individual MongoDB Pods**.

An application connecting to MongoDB can reach specific replicas like:

```
mongo --host mongo-0.mongo-service
```

## ◆ When Should You Use Headless Services?

- ◆ **Databases (MongoDB, MySQL, PostgreSQL, Cassandra)** - When clients need to connect to specific instances.

- ◆ **Message Brokers (Kafka, RabbitMQ)** - Where direct communication with nodes is required.

- ◆ **Peer-to-Peer Applications** - Services like Elasticsearch that require direct communication.

- ◆ **Custom Load Balancing** - If an external load balancer (like Envoy or Nginx) manages traffic instead of Kubernetes.

---

## Real-World Example: Headless Service for a MongoDB Replica Set

Imagine you're deploying **MongoDB** in Kubernetes as a **Replica Set**. This setup ensures **high availability and failover** for your database.

### ◆ Scenario

You're running a **three-node MongoDB cluster** on Kubernetes using a **StatefulSet**. Your application needs to connect to each MongoDB Pod **individually** instead of a load-balanced service.

### 1 Define a Headless Service for MongoDB

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: mongo-headless
```

```
spec:
```

---

```
clusterIP: None # Makes it headless (no single IP)
```

```
selector:
```

```
  app: mongo
```

```
ports:
```

```
  - port: 27017
```

```
    targetPort: 27017
```

♦ **Why is this needed?**

- Since **clusterIP: None**, it won't provide a single IP.
- Instead, Kubernetes **DNS will return all individual Pod IPs**.

## 2 Create a MongoDB StatefulSet

```
apiVersion: apps/v1
```

```
kind: StatefulSet
```

```
metadata:
```

---

```
name: mongo
```

```
spec:
```

```
  serviceName: mongo-headless # Connects to the Headless
  Service
```

```
  replicas: 3 # Three MongoDB replicas
```

```
  selector:
```

```
    matchLabels:
```

```
      app: mongo
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: mongo
```

```
    spec:
```

```
      containers:
```

```
        - name: mongo
```

```
          image: mongo:latest
```

---

```
    ports:

    - containerPort: 27017

    volumeMounts:

    - name: mongo-storage

      mountPath: /data/db

  volumeClaimTemplates:

  - metadata:

      name: mongo-storage

    spec:

      accessModes: ["ReadWriteOnce"]

      resources:

        requests:

          storage: 1Gi
```

#### ♦ How Does This Work?

- Kubernetes **assigns stable DNS names** to each Pod:



- 
- `mongo-0.mongo-headless.default.svc.cluster.local`
  - `mongo-1.mongo-headless.default.svc.cluster.local`
  - `mongo-2.mongo-headless.default.svc.cluster.local`
- The **MongoDB replica set** can now communicate using these DNS names.

### 3 Initialize MongoDB Replica Set

Once the StatefulSet is running, you need to initialize the **MongoDB replica set** inside one of the Pods.

#### Step 1: Enter the First MongoDB Pod

```
kubectl exec -it mongo-0 -- mongo
```

#### Step 2: Initialize the Replica Set

```
rs.initiate({  
  
  _id: "rs0",  
  
  members: [  

```

---

```
    { _id: 0, host:
"mongo-0.mongo-headless.default.svc.cluster.local:27017" },

    { _id: 1, host:
"mongo-1.mongo-headless.default.svc.cluster.local:27017" },

    { _id: 2, host:
"mongo-2.mongo-headless.default.svc.cluster.local:27017" }

]

}))
```

♦ **What This Does:**

- It **sets up the replica set** using the individual Pod hostnames.
- Each MongoDB instance **knows the other instances** and can sync data.

#### **4 Application Connection**

Your application can now connect to the MongoDB **replica set** using:

```
mongodb://mongo-0.mongo-headless.default.svc.cluster.local:27017,mongo-1.mongo-headless.default.svc.cluster.local:27017,
```

---

```
mongo-2.mongo-headless.default.svc.cluster.local:27017/?replicaSet=rs0
```

#### ♦ Why is This Important?

- The app connects to **all MongoDB instances**.
- If a **Pod goes down**, MongoDB will elect a **new primary node** automatically.
- The system remains **highly available**.

## Why Not Use a Normal (ClusterIP) Service for Databases Like MongoDB?

A **normal ClusterIP service** provides a **single IP** that load balances traffic across multiple Pods. While this works for **stateless applications**, it's problematic for **stateful databases like MongoDB, Cassandra, or Kafka**. Here's why:

#### ♦ 1 Stateful Databases Require Direct Pod Communication

- A database like **MongoDB (Replica Set)** or **Kafka (Broker Nodes)** needs each instance to **communicate directly** with others.

- 
- If we use a **ClusterIP Service**, the traffic gets **randomly load-balanced** across all Pods.
  - This **breaks** leader elections, replication, and node-to-node communication.

#### Example Issue with a ClusterIP Service:

```
mongodb://my-service.default.svc.cluster.local:27017
```

- The service forwards requests **randomly** to `mongo-0`, `mongo-1`, or `mongo-2`.
- MongoDB **needs to connect to specific nodes**, but it **can't distinguish** them.
- This **prevents proper Replica Set initialization**.

#### ♦ 2 Primary and Replica Nodes Need a Stable Identity

- In a **MongoDB Replica Set**, one node is elected as **Primary**, and others as **Replicas**.
- If a **Replica wants to sync data**, it must know the **exact hostname** of the Primary.

- 
- A **ClusterIP Service** hides the **Pod IPs**, making it **impossible to form a proper Replica Set**.

**Example:**

 **With Headless Service (Correct)**

```
mongodb://mongo-0.mongo-headless.default.svc.cluster.local:27017,mongo-1.mongo-headless.default.svc.cluster.local:27017,mongo-2.mongo-headless.default.svc.cluster.local:27017/?replicaSet=rs0
```

- Each MongoDB instance can **directly talk** to others using **Pod DNS names**.
- If **mongo-0** becomes Primary, replicas know exactly where to send write requests.

 **With ClusterIP Service (Incorrect)**

```
mongodb://mongo-service.default.svc.cluster.local:27017
```

- The **service randomly routes traffic**, so Replicas can't find the **Primary node**.

- 
- MongoDB **fails to maintain replication** and **loses consistency**.

### ◆ 3 Stateful Workloads Need Persistent Storage

- ClusterIP services work well for **stateless applications** like APIs, but databases need **stable storage**.
- **StatefulSets + Headless Services** ensure: ☒ Each database Pod has a **stable identity** (like `mongo-0`, `mongo-1`).  
☒ Each Pod gets its own **PersistentVolume** for storing data.

### ◆ 4 ClusterIP Service Breaks Database Failover

- If the **Primary node dies**, MongoDB must **elect a new leader**.
- With a **Headless Service**, all replicas **immediately update their connections** to the new leader.
- With a **ClusterIP Service**, the app **won't know which node is the new Primary**.


---

# How DNS Works in Kubernetes Services?

Kubernetes provides **built-in DNS** to allow services, pods, and other resources to communicate using domain names instead of IP addresses. This is essential because **Pods are dynamic**, and their IP addresses keep changing.

## 1 How Does Kubernetes DNS Work?

Kubernetes uses **CoreDNS** as its built-in DNS service. It automatically assigns a **DNS name** to each Service and Pod inside the cluster.

 **When a Service is created**, Kubernetes **registers it in the internal DNS**, allowing other Pods to access it using a domain name.

### DNS Naming Format for Services

`<ServiceName>.<Namespace>.svc.cluster.local`

- **ServiceName** → The name of the Service.
- **Namespace** → The Kubernetes namespace where the Service is deployed.
- **svc.cluster.local** → The default cluster domain.

---

## 2 Example of DNS Resolution in Kubernetes

### ◆ Scenario:

- You have a **backend service** named `backend-svc` running in the `default` namespace.
- A **frontend pod** wants to communicate with it.

### ✅ DNS Resolution in Action

The frontend pod can use the **service name** directly to connect:

`http://backend-svc`

If the pod is in a different namespace, it must use:

`http://backend-svc.default.svc.cluster.local`

- CoreDNS resolves this to the **ClusterIP** of `backend-svc`, forwarding requests to the correct pods.

## 3 DNS Resolution for Headless Services

### ◆ Regular Services (ClusterIP, NodePort, LoadBalancer):



- 
- Resolve to a **single ClusterIP**, which load balances traffic to multiple pods.

- ◆ **Headless Services (`clusterIP: None`):**

- Instead of a **single IP**, CoreDNS returns a list of **Pod IPs**.
- This allows direct pod-to-pod communication, useful for databases like MongoDB or Cassandra.

📌 **Example:** If `mongo-service` is headless (`clusterIP: None`), running:

```
nslookup mongo-service.default.svc.cluster.local
```

will return a list of **MongoDB pod IPs** instead of a single ClusterIP.

#### 4 How DNS Resolves Service Names?

When a pod makes a request to `backend-svc`, Kubernetes DNS resolves it in **three steps**:

1 **CoreDNS intercepts the request** and checks its records.

2 **If a match is found**, it returns the ClusterIP (for normal services) or Pod IPs (for headless services).

---

③ The request is then forwarded to the correct **pod** based on the Service's selector.

## What Happens Behind the Scenes When a Kubernetes Service is Created?

When you create a Kubernetes Service, a lot of things happen in the background to ensure network connectivity, load balancing, and service discovery. Let's break it down step by step.

### ① Step 1: Service Definition & API Request

When you apply a Service YAML like this:

```
apiVersion: v1

kind: Service

metadata:

  name: my-service

spec:
```

---

```
selector:
```

```
  app: my-app
```

```
ports:
```

```
  - protocol: TCP
```

```
    port: 80
```

```
    targetPort: 8080
```

- **Kubernetes API Server** receives the request.
- It stores the Service definition in **etcd (the Kubernetes key-value store)**.
- The Service is now registered in the cluster.

## **2** Step 2: Service Gets a ClusterIP (if applicable)

- If the Service type is **ClusterIP (default)**, Kubernetes assigns a **virtual IP (ClusterIP)**.
- This **IP never changes** as long as the Service exists.
- The ClusterIP is stored in **kube-proxy** running on each node.

---

### 3 Step 3: EndpointSlice Controller Finds Matching Pods

- Kubernetes checks **which Pods match the selector** (**app: my-app**).
- It creates an **EndpointSlice** that stores the IPs of all matched pods.

Example of an EndpointSlice:

```
apiVersion: discovery.k8s.io/v1
```

```
kind: EndpointSlice
```

```
metadata:
```

```
  name: my-service-abc123
```

```
addressType: IPv4
```

```
endpoints:
```

```
  - addresses:
```

```
    - 10.244.1.12 # Pod 1 IP
```

```
    - 10.244.2.25 # Pod 2 IP
```

```
ports:
```

---

- name: http

- port: 8080

- protocol: TCP

- 

- Kube-proxy on each node watches for these EndpointSlices.

## 4 Step 4: Kube-Proxy Updates IPTables/IPVS Rules

- **kube-proxy** (running on each node) updates **iptables** or **IPVS** rules to forward traffic.
- If using **iptables**, a rule is added to forward requests to one of the Pod IPs in a round-robin way.
- If using **IPVS**, it creates a **load balancing table**.

- ♦ **Example of iptables rules created:**

```
-A KUBE-SERVICES -d 10.96.0.10 -p tcp --dport 80 -j
```

```
KUBE-SVC-XXXXXXX
```

```
-A KUBE-SVC-XXXXXXX -m statistic --mode random --probability
```

```
0.5 -j KUBE-SEP-YYYYYYY
```

---

```
-A KUBE-SEP-YYYYYYYY -s 10.244.1.12 -j DNAT --to-destination  
10.244.1.12:8080
```

- Requests to **ClusterIP (10.96.0.10:80)** are forwarded to **Pod IPs (10.244.1.12:8080, etc.)**.

## 5 Step 5: CoreDNS Updates DNS Records

- **CoreDNS**, the Kubernetes DNS system, gets updated.

A new DNS entry is added:

```
my-service.default.svc.cluster.local → 10.96.0.10  
(ClusterIP)
```

- 

- Now, any pod can resolve `http://my-service` to its **ClusterIP**.

---

## 6 Step 6: Traffic Flowing to Service

### What Happens When a Pod Sends a Request to the Service?

A pod tries to access the service:

```
curl http://my-service
```

1. DNS resolves `my-service` to **ClusterIP (10.96.0.10)**.
2. **Kube-proxy intercepts** the request and **forwards it to a Pod IP** (e.g., 10.244.1.12:8080).
3. The target **Pod receives the request and processes it**.

## 7 What Happens When Pods Scale Up or Down?

- The **EndpointSlice Controller** constantly monitors Pods.
- If new Pods are added (e.g., autoscaling), their IPs are automatically added to the EndpointSlice.
- If Pods are deleted, their IPs are removed.
- `kube-proxy` updates its rules accordingly.

---

## 8 What Happens If a Service is Deleted?

- **The ClusterIP is removed** from kube-proxy.
- **EndpointSlice is deleted**, so Pods are no longer linked to the Service.
- **DNS records are removed**, so the Service name won't resolve anymore.

## Does Kubernetes Create EndpointSlices for Each Service?

Yes, **Kubernetes automatically creates EndpointSlices** for each Service that has a selector. The **EndpointSlice controller** manages them and updates them dynamically as Pods are added or removed.

### ♦ How EndpointSlices Work

- When you create a Service, Kubernetes automatically creates an **EndpointSlice** containing all matching Pod IPs.
- If a Pod is deleted, Kubernetes **removes its IP** from the EndpointSlice.
- If new Pods matching the selector come up, Kubernetes **adds their IPs** to the EndpointSlice.



---

## What Happens If You Manually Remove a Pod IP from EndpointSlice?

If you manually edit the EndpointSlice and remove a Pod's IP:

1. **The Pod will still be running**, but the Service will no longer route traffic to it.
2. **Kube-proxy will not forward requests** to that Pod anymore.
3. The Pod will still be accessible **directly via its Pod IP** but **not via the Service**.

💡 **However, Kubernetes will automatically restore the removed Pod IP** because the EndpointSlice controller continuously syncs it. So, the change is temporary.

## What Happens If You Manually Add a Pod IP to EndpointSlice?

If you manually add an IP of a non-existing or unrelated Pod:

1. **Kube-proxy will try to send traffic** to that IP.
2. If the IP does not belong to an active Pod, requests will **fail or be misrouted**.
3. Kubernetes **may override your changes** because the EndpointSlice controller continuously updates the list based on actual Pods.

💡 **Manual modifications to EndpointSlices are not recommended** because Kubernetes will eventually correct the changes.

---

## How to Control Which Pods Get Added to the EndpointSlice?

- Instead of modifying EndpointSlices manually, **use proper Service selectors**.
- If you want **manual control over endpoints**, use a **Headless Service (ClusterIP: None)** and manage endpoints yourself.

## ◆ Summarizing Kubernetes Services ◆

---

Kubernetes Services play a **crucial role** in enabling seamless communication between Pods, providing a stable endpoint despite the **dynamic and ever-changing** nature of Kubernetes environments. They abstract the underlying complexity of networking and ensure that applications remain **discoverable, scalable, and resilient**. 🚀

From **ClusterIP for internal communication** to **NodePort and LoadBalancer for external access**, Services offer various options to fit different deployment needs. Additionally, **Headless Services** and **ExternalName Services** provide even more flexibility for handling DNS-based routing and custom networking setups. 🌐

By leveraging **Selectors, Endpoints, and Service Discovery mechanisms**, Kubernetes Services ensure that traffic is efficiently routed to healthy Pods, balancing the load and maintaining high availability. **Integrating them with Ingress and DNS further enhances their power**, making Kubernetes a truly production-ready orchestration system. 💡

As we wrap up this discussion, remember that **understanding and effectively using Services is key to managing Kubernetes networking efficiently**. Keep practicing, experimenting, and deep-diving into Kubernetes – the more you

---

explore, the more powerful your infrastructure becomes! 💪

◆ **Wrapping Up & Looking Ahead** ◆

---

As we conclude this deep dive into **Kubernetes Services**, I want to take a moment to express my gratitude for your **continuous support and engagement** in this Kubernetes learning journey. 🙌 Your enthusiasm and curiosity drive me to share more valuable insights, ensuring that we all grow together in mastering Kubernetes. 🚀

Kubernetes Services are **a fundamental component of cluster networking**, ensuring seamless communication between Pods and enabling efficient traffic routing. Mastering them is key to building **scalable, resilient, and production-ready applications** in Kubernetes. 🌐

But this is just the beginning! **Stay tuned for more hands-on Kubernetes tasks**, where we'll continue exploring essential concepts, practical implementations, and real-world use cases. 💡

👉 **Follow me for more daily Kubernetes content** filled with practical examples, best practices, and deep insights. Let's keep learning, building, and growing together! 💪🔥

♦ **Thank you for being part of this journey!** Keep supporting, keep learning, and stay connected for more. 🚀