

AWS SERVICES REAL TIME DAILY TASKS

AWS SERVERLESS ARCHITECTURE

COMPREHENSIVE GUIDE

THEORY + PRACTICAL

WITH REAL TIME PROJECT

FOLLOW FOR MORE

FENILGAJJAR.DEVOPS@GMAIL.COM



Serverless Architecture:

with Theory + Real-Time

Project Implementation!



✓ Welcome to the Ultimate Guide on AWS Serverless Architecture!

In today's evolving cloud landscape, **serverless architecture** is revolutionizing the way applications are built, deployed, and scaled. Gone are the days of managing complex server infrastructures—AWS Serverless enables developers to focus on code while AWS takes care of scalability, availability, and maintenance!

My What You'll Learn in This Guide:

- **Deep Dive into Serverless Architecture** Understand its core principles, benefits, and real-world use cases.
- AWS Serverless Services Explore AWS Lambda, API Gateway, DynamoDB, S3, Step Functions, and more.
- ✓ Hands-on Project Implementation Step-by-step walkthrough of building a real-world serverless application from scratch.

Best Practices & Optimization – Learn how to enhance performance, reduce costs, and improve security in serverless environments.

By the end of this guide, you'll gain **both theoretical knowledge and practical experience** to confidently design and implement serverless solutions on AWS!

** Stay tuned for the complete walkthrough, and let's unlock the full potential of AWS Serverless Architecture together!



What is AWS Serverless Architecture?

AWS Serverless Architecture is a cloud computing model that allows developers to build and run applications without managing servers. Instead of provisioning and maintaining infrastructure, AWS handles everything behind the scenes, including scalability, availability, and security.

With serverless architecture, you focus solely on writing code, and AWS takes care of executing it in response to events. You are **charged only for the exact execution time** rather than paying for always-on servers.

Key Components of AWS Serverless Architecture

_1 AWS Lambda (Compute Service)

AWS Lambda is the core of serverless computing in AWS. It lets you run code without provisioning or managing servers.

- Event-driven execution Lambda runs in response to triggers like API requests, file uploads, or database changes.
- Supports multiple languages Python, Node.js, Java, Go, and more.
- Auto-scaling AWS automatically scales your function based on demand.
- Pay-per-use pricing You pay only for the execution time.
- * Example: A Lambda function can trigger when a user uploads an image to an S3 bucket, automatically resizing the image.

2 Amazon API Gateway (Serverless API Management)

Amazon API Gateway is a fully managed service that enables you to **create**, **deploy**, **and manage RESTful and WebSocket APIs**.

- Connects frontend apps to AWS Lambda to execute business logic.
- Supports authentication & authorization with AWS IAM, Cognito, and JWT tokens.
- Handles traffic spikes without manual intervention.
- Integrates with AWS WAF to protect APIs from malicious attacks.

* Example: A mobile app makes an HTTP request to an API Gateway, which triggers an AWS Lambda function to fetch data from DynamoDB.

3 Amazon DynamoDB (Serverless NoSQL Database)

Amazon DynamoDB is a **serverless NoSQL database** that provides **fast and scalable performance** for applications.

- Fully managed No need to handle database servers.
- Supports auto-scaling Dynamically adjusts read/write capacity.
- Highly available & fault-tolerant Replicates data across multiple regions.
- Integrated with AWS Lambda Lambda can fetch or update DynamoDB records.

* Example: A real-time leaderboard app stores player scores in DynamoDB, and AWS Lambda updates scores when a game ends.

4 Amazon S3 (Serverless Storage)

Amazon S3 (Simple Storage Service) is an **infinitely scalable** object storage solution in AWS.

- Stores unstructured data Images, videos, logs, backups, and static files.
- Highly durable & available 99.99999999% (11 nines) of durability.
- Integrates with AWS Lambda Triggers functions on file uploads.
- **Lifecycle policies** Automatically move files to cheaper storage classes like Glacier.

* Example: A website stores static images on S3, which are served via Amazon CloudFront for fast delivery.

5 Amazon EventBridge & SNS (Event-Driven Messaging)

AWS EventBridge and SNS (Simple Notification Service) provide event-driven communication between AWS services.

- **EventBridge** routes events between AWS services and third-party SaaS applications.
- SNS enables real-time notifications via SMS, email, or push notifications.
- Ensures decoupling Microservices can communicate asynchronously.

* Example: An e-commerce site sends an order confirmation email using SNS when a customer places an order.

6 AWS Step Functions (Serverless Workflow Orchestration)

AWS Step Functions allow developers to **orchestrate multiple AWS services into** a single workflow.

- Manages complex workflows using state machines.
- Integrates with Lambda, DynamoDB, S3, and more.
- Visual workflow designer for easy debugging and monitoring.
- **Reduces application complexity** by eliminating the need for custom workflow logic.

* Example: A video processing pipeline automatically transcodes uploaded videos using Step Functions and AWS Lambda.

Why Should You Use AWS Serverless Architecture?



AWS Serverless Architecture offers **scalability**, **cost efficiency**, **and operational simplicity**, making it an ideal choice for modern applications. Below are the key reasons why you should use it:

🔟 No Server Management 🎇

With serverless, you don't need to **provision, configure, or maintain servers**. AWS fully manages the infrastructure, including:

- ✓ Automatic scaling ✓
- ✓ High availability & fault tolerance ✓

* Example: Instead of setting up and maintaining an EC2 instance to process API requests, you can use **AWS Lambda + API Gateway**, which auto-scales without any manual intervention.

2 Cost Efficiency 💰

AWS Serverless follows a pay-as-you-go model, meaning you are billed only for the actual execution time rather than 24/7 server uptime.

✓ No cost when your functions are idle.

- ✓ No need to over-provision resources for peak loads.
- ✓ Millisecond-based billing in AWS Lambda.

* Example: A traditional EC2 instance runs 24/7 and incurs costs even when idle, whereas AWS Lambda runs only when triggered, significantly reducing costs.

Automatic Scaling \(\psi \)

Serverless services scale automatically based on demand.

- ✓ AWS Lambda can handle a few requests per second or millions per minute.
- ✓ Amazon DynamoDB auto-scales based on traffic without downtime.
- ✓ API Gateway manages sudden traffic spikes efficiently.

* Example: A flash sale e-commerce website can seamlessly scale during peak hours without any downtime using AWS Serverless.

4 Faster Development & Deployment 🧲

AWS Serverless eliminates infrastructure setup, allowing developers to:

- ✓ Focus on writing code instead of managing servers.
- ✓ Deploy updates instantly without server restarts.
- ✓ Use microservices for modular and scalable applications.
- * Example: A startup building a real-time chat app can quickly launch using Lambda, API Gateway, and DynamoDB, reducing time-to-market.

5 Built-in Fault Tolerance & High Availability 🔽



AWS automatically replicates serverless functions across multiple Availability Zones (AZs) for:

- √ 99.999% uptime (high availability).
- ✓ Automatic failover handling in case of failures.
- ✓ Data durability with AWS S3, DynamoDB, and RDS.

* Example: A serverless video streaming platform remains operational even if one AWS data center fails.

6 Security & Compliance 🔐

AWS handles security updates, patches, and compliance requirements for serverless services:

- ✓ AWS Lambda runs in isolated execution environments (sandboxed).
- ✓ IAM roles restrict access to AWS resources.
- **✓** API Gateway provides authentication & rate limiting.

* Example: A healthcare app using AWS Serverless meets HIPAA compliance without additional security setup.

$oldsymbol{7}$ Event-Driven & Real-Time Processing $oldsymbol{ eq}$

Serverless architecture is perfect for **event-driven applications** like:

✓ Real-time data processing (e.g., log monitoring, analytics).

- ✓ File processing (e.g., image resizing, PDF generation).
- ✓ **IoT applications** (e.g., processing sensor data).
- * Example: AWS Lambda can trigger whenever a file is uploaded to S3, automatically converting an image to different formats.

8 Best for Microservices & Modern Apps 🟗

AWS Serverless aligns with microservices architecture, enabling:

- ✓ Independent deployments without affecting other services.
- ✓ Decoupled services using event-driven communication (SNS, SQS, EventBridge).
- ✓ Easier maintenance & scalability.
- * Example: A social media app using separate Lambda functions for user authentication, notifications, and media uploads makes scaling easier.

Evolution of Cloud Computing: Before and After AWS Serverless Architecture

AWS Serverless didn't exist from the beginning—cloud computing evolved over time. Let's break it down into what existed before AWS Serverless, how Serverless transformed cloud computing, and what's next after Serverless.

Before AWS Serverless Architecture (Traditional & Cloud-Based Approaches)

1 On-Premises Infrastructure (Pre-Cloud Era)

- Companies bought physical servers, managed hardware, storage, and networking.
- High cost (buying & maintaining servers, electricity, cooling).
- Limited scalability (had to predict future needs).
- Downtime risks (if a server failed, everything could go down).

* Example: A bank running an in-house data center with dedicated servers for different applications.

2 Virtual Machines (VMs) & Private Cloud

- Companies virtualized their servers using technologies like VMware.
- Enabled better resource utilization but still required manual maintenance.
- Improved scalability compared to physical servers but not automatic.

* Example: Instead of running 10 physical servers, a company ran multiple VMs on a few powerful servers.

③Infrastructure as a Service (IaaS) – AWS EC2 (Early Cloud) 🧢

- Cloud providers (AWS, Azure, GCP) introduced virtual machines (EC2 instances).
- Businesses rented servers on demand instead of buying hardware.
- Elastic scaling (could add/remove VMs but required configuration).
- Still required server maintenance (OS updates, scaling, security patches).

* Example: A startup using AWS EC2 instances to host a website instead of buying physical servers.

4 Platform as a Service (PaaS) – AWS Elastic Beanstalk 🌱

- AWS introduced PaaS solutions to reduce server management.
- Developers deployed applications without worrying about OS, networking,
 and scaling.
- AWS handled provisioning, load balancing, and scaling automatically.
- However, still required managing runtime environments.

* Example: A developer deploying a Python app on Elastic Beanstalk, which automatically set up EC2, load balancers, and databases.

* AWS Serverless Architecture Changes Everything! (2014

- Present) 🚀

Introduction of Serverless (AWS Lambda & Beyond)

AWS **Lambda (2014)** introduced **serverless computing**, where developers run code without managing infrastructure.

- **✓ No server management** AWS handles provisioning, scaling, and patching.
- ✓ Pay-per-use model Costs only for execution time (no idle charges).
- ✓ Event-driven execution Runs in response to triggers (S3, API Gateway, DynamoDB, etc.).
- **✓ Microservices-friendly** Functions scale independently.
- * Example: A user uploads an image to **S3**, triggering **AWS Lambda** to resize it and store the resized version in another S3 bucket.

Serverless components include:

- AWS Lambda (Compute)
- API Gateway (Managed API)
- DynamoDB (Serverless Database)
- **\$3** (Serverless Storage)
- Step Functions (Serverless Workflows)

What's After AWS Serverless? (The Future) 🚀

- Hybrid of serverless and containerization.
- Runs Docker containers without managing EC2 instances.
- Ideal for long-running processes that AWS Lambda doesn't support.

P Example: Running a machine learning model in AWS Fargate instead of provisioning an EC2 instance.

2 Edge Computing (AWS Lambda@Edge & AWS Wavelength)



- Moving computation closer to users (on CloudFront edge locations).
- Reduces latency for real-time applications (IoT, gaming, 5G, Al).

* Example: A global video streaming platform using Lambda@Edge to transcode videos based on user location.

3Al-Driven Serverless Computing 🤖

- Future serverless will auto-optimize resources using Al.
- Services like AWS Bedrock & SageMaker will offer serverless AI model hosting.

* Example: A chatbot automatically scales based on user queries using Al-driven serverless services.

When Should You Use AWS Serverless

Architecture?

AWS Serverless is **not a one-size-fits-all solution**—it's powerful but works best for specific use cases. Here's when you should consider using AWS Serverless:

✓ 1 Event-Driven Applications

When?

- ✓ Your application needs to react to events (file uploads, API requests, database changes).
- ✓ You don't need a continuously running server.

Example:

- An image processing pipeline where users upload images to S3, triggering
 AWS Lambda to resize them.
- A real-time notification system where DynamoDB streams trigger Lambda to send alerts.

✓ 2 Microservices & API-Driven Applications <a> □

When?

- ✓ You are building small, independent services that communicate via APIs.
- ✓ You need auto-scaling APIs without managing backend servers.

📌 Example:

- A REST API powered by AWS Lambda + API Gateway without managing an EC2 server.
- A GraphQL backend using AWS AppSync and DynamoDB.

✓ 3 Unpredictable Traffic & Auto-Scaling Needs

When?

- ✓ Your app has spiky or variable traffic (some days high, some days low).
- ✓ You want instant auto-scaling without provisioning servers.

Example:

- A **ticket booking system** that gets high traffic only during special events.
- An **IoT backend** processing thousands of device messages in real time.

When?

- ✓ You want to pay only for what you use instead of always running a server.
- ✓ Your app has low-to-moderate usage (serverless is expensive for heavy, always-on workloads).

Example:

- A **chatbot** that only runs when a user sends a message.
- A cron job replacement where AWS Lambda runs once per hour instead of keeping an EC2 instance running.

☑ 5 Fast Prototyping & MVP Development

When?

- ✓ You need to quickly develop a proof of concept (POC) or Minimum Viable Product (MVP).
- ✓ You don't want to spend time on server setup and infrastructure management.

P Example:

- A startup building an app prototype using AWS Lambda + DynamoDB + API Gateway.
- A hackathon project that needs fast deployment.

V6 Serverless Data Processing & Automation ☐

When?

- ✓ You need to process data in real-time or batch mode without running a server.
- ✓ Your workflow is triggered by an **event (file upload, database update, message queue, etc.).**

Example:

- **Log processing** where CloudWatch Logs trigger AWS Lambda for real-time analysis.
- ETL pipelines where AWS Glue and Step Functions automate data transformation.

X When NOT to Use AWS Serverless?

- X Long-Running Applications (24/7 Services)
- AWS Lambda has a maximum execution time of 15 minutes.
- Better alternative: Use AWS Fargate (for containers) or EC2 instances for long-running tasks.

Example:

• Running a video streaming server or a real-time multiplayer game server.

X High-Performance Computing (HPC) & Low-Latency Workloads 🚀

- Serverless adds cold start latency (when functions start after inactivity).
- Better alternative: Use EC2 or AWS Batch for HPC.

Example:

Machine learning training (use SageMaker or EC2 GPU instances instead).

X Large Monolithic Applications îm

- Serverless works best for microservices and event-driven apps.
- Better alternative: Use EC2, Kubernetes (EKS), or Elastic Beanstalk for monolithic applications.

P Example:

• Legacy enterprise applications with stateful architecture.

How Does AWS Serverless Architecture Work?

AWS Serverless Architecture removes the need for manual server management. Instead, AWS automatically provisions, scales, and manages the infrastructure. Let's break down how it works step by step.

Key Components of AWS Serverless Architecture

🔟 Event-Driven Execution 🚀

- Serverless applications only run when triggered by an event (e.g., an API request, file upload, database update).
- No need to keep servers running all the time.

2 Auto-Scaling & Pay-Per-Use 💰

- AWS scales your application automatically based on demand.
- You pay only for what you use, not for idle servers.

- AWS handles provisioning, maintenance, scaling, and availability of resources.
- No need to manage or configure servers manually.

Step-by-Step Workflow of AWS Serverless Architecture

Step 1: An Event Triggers the Application 📩

AWS Serverless applications are **event-driven**—they start running **only when an event occurs**.

Common Event Sources:

- API Request: A user requests data from a web or mobile app.
- File Upload: A file is uploaded to an S3 bucket.
- **Database Change:** A new record is inserted into DynamoDB.
- Message Queue: A message arrives in an SQS queue.

Example: A user uploads an image to S3. This event triggers an AWS Lambda function to resize the image.

Step 2: AWS Lambda Executes the Business Logic 🔅

Once an event occurs, **AWS Lambda** (or another serverless compute service) runs the required function.

How Lambda Works?

- AWS automatically allocates resources and runs the function.
- The function executes within milliseconds and automatically scales up/down.
- Once execution is complete, AWS **deallocates the resources** to save costs.



 Lambda fetches the uploaded image, resizes it, and saves the processed image back to S3.

Step 3: AWS Services Handle Data Storage & API Requests III

Since serverless apps don't use traditional servers, they **leverage AWS-managed** services for data storage and API handling.

- Common AWS Services Used in Serverless Apps:
 - **DynamoDB** → NoSQL database for storing app data.
 - **S3** → Object storage for files & images.
 - API Gateway → Exposes APIs for frontend apps.
 - **Step Functions** → Manages workflows between different Lambda functions.

Example:

- After resizing an image, Lambda stores it in an S3 bucket.
- API Gateway allows a web app to fetch the resized image.

Step 4: AWS Automatically Scales the Application 📈

One of the biggest advantages of serverless is **auto-scaling**.

- How Scaling Works?
 - If **10 users** upload images, AWS Lambda automatically **creates 10 instances** of the function.
 - If **100,000 users** upload images, AWS scales instantly to handle the load.
 - No manual intervention is needed—AWS manages everything.

Example:

 A viral social media app gets 1 million API requests in 1 hour. AWS automatically scales API Gateway & Lambda to handle it.

Step 5: The Application Runs Only When Needed, Reducing Costs 💰



Unlike traditional servers that run 24/7, AWS Serverless only runs when triggered.

- Cost Optimization in Serverless:
 - You only pay for the compute time used by Lambda (measured in milliseconds).
 - No need to pay for idle servers.
 - AWS deallocates resources once the function execution is complete.

Example:

• If a chatbot function runs for 2 seconds per request, you only pay for those 2 seconds—not for an entire server.

Real-World Example: How AWS Serverless Works in a Web App

Use Case: A serverless e-commerce website for ordering products.

Step 1: User Places an Order

User clicks "Buy Now" on a website.

• API Gateway sends the request to AWS Lambda.

** Step 2: AWS Lambda Processes the Order

- Lambda checks inventory in DynamoDB.
- Lambda calculates the total price and saves order details.

Step 3: Payment Processing

- Lambda calls Stripe API to process payment.
- If payment is successful, it stores order details in DynamoDB.

Step 4: Order Confirmation Email Sent

- AWS SNS (Simple Notification Service) sends an email to the customer.
- The order details are stored in an S3 bucket for analytics.

Poutcome: The entire process happens without managing servers, and the app automatically scales during high traffic.



Popular Use Case Scenarios of AWS Serverless

Architecture

AWS Serverless Architecture is widely used across different industries to build cost-effective, scalable, and event-driven applications. Here are some real-world scenarios where AWS Serverless fits perfectly:

Web & Mobile Backend Development

- Serverless is great for building scalable web and mobile backends.
- AWS Lambda processes API requests, Amazon API Gateway manages endpoints, and **DynamoDB** or **Aurora Serverless** stores data.

P Example:

- A social media app stores user data in DynamoDB and serves profile images from S3.
- A chat application processes messages via WebSockets using API Gateway and Lambda.

Real-Time Image & Video Processing

• Serverless is perfect for automating image and video processing without managing infrastructure.

 When an image/video is uploaded to Amazon S3, it triggers a Lambda function to process it.

P Example:

- A photo-sharing app automatically resizes images and generates thumbnails when users upload photos.
- A video streaming platform converts uploaded videos into multiple formats for different devices.

E-Commerce & Order Processing

- E-commerce applications need fast, scalable order processing with minimal infrastructure management.
- Serverless can handle user authentication, payment processing, and real-time inventory management.

Example:

- A serverless shopping cart system updates stock availability in DynamoDB and processes orders with Lambda.
- Payments are handled using Stripe API, and order confirmation emails are sent via SNS (Simple Notification Service).

III Data Processing & ETL (Extract, Transform, Load)

 AWS Serverless is widely used in big data pipelines for processing large datasets in real time. Lambda, Kinesis, Glue, and Athena help ingest, transform, and analyze data.

P Example:

- A finance company processes real-time stock price updates and alerts users on significant price changes.
- A marketing analytics platform collects user behavior data from multiple sources and stores it in an S3-based data lake.

IoT (Internet of Things) Applications

- AWS Serverless is ideal for handling millions of IoT device connections efficiently.
- AWS IoT Core handles device communication, and Lambda processes incoming data in real-time.

Example:

- A smart home automation system automatically adjusts lights and temperature based on sensor data.
- A fleet tracking system collects real-time vehicle location data and updates it in DynamoDB.

Security & Authentication Systems

• Serverless can securely manage user authentication and access control.

 Amazon Cognito provides user authentication, and Lambda customizes authentication workflows.

P Example:

- A banking app authenticates users via Cognito and validates transactions using Lambda functions.
- A gaming platform uses Cognito for login and API Gateway to manage secure API access.

Serverless Chatbots & Virtual Assistants

- AWS serverless is commonly used to build chatbots and Al-powered assistants.
- Amazon Lex, Lambda, and DynamoDB power conversational Al applications.

Example:

- A customer support chatbot automates responses for a company's website and WhatsApp support.
- A voice assistant integrates with Alexa using AWS Lambda for smart home control.

Serverless Machine Learning Inference

 Serverless allows running ML models at scale without provisioning servers. SageMaker, Lambda, and Step Functions work together for machine learning inference.

P Example:

- A fraud detection system in banking analyzes transactions in real-time for anomalies.
- A personalized recommendation engine suggests products based on user behavior.

Global Content Delivery & Static Websites

- AWS Serverless makes it easy to host static websites with automatic scaling and low cost.
- S3, CloudFront, and API Gateway deliver content globally with minimal latency.

Example:

- A company's blog is hosted on S3 and distributed via CloudFront with API
 Gateway for backend services.
- A portfolio website is deployed as a fully serverless site using S3 and Lambda@Edge for dynamic content.

Welcome to the Real-Time Implementation of Serverless Architecture with DynamoDB, API Gateway, and Lambda!

In this project, we are taking a **deep dive into AWS Serverless Architecture**, where we will build a **fully functional**, **scalable**, **and event-driven backend system**—without managing any servers!

- No Infrastructure Management Focus on code, not servers.
- Lightning-Fast Data Processing Leverage DynamoDB for high-speed,
 low-latency storage.
- Scalable API Endpoints Use API Gateway to expose serverless functions securely.
- Event-Driven Execution Let AWS Lambda handle the backend logic dynamically.

★ What's Inside?

- A step-by-step real-world implementation of a fully serverless bookstore application.
- Seamless API-driven workflow where API requests trigger Lambda to interact with DynamoDB.
- **Testing, debugging, and optimization techniques** for a robust serverless solution.
- Best practices for securing and optimizing AWS Serverless workflows.

This guide will cover both the theoretical concepts and a hands-on real-time project, ensuring you get a solid grasp of AWS Serverless Services and their real-world applications.

Get ready to build, automate, and scale without managing a single server!

Stay tuned for a detailed walkthrough and implementation!

Real Time Task on Step By Step

Implementation of Serverless Architecture

with DynamoDB, Api Gateway and Lambda.

Step 1: Create a DynamoDB Table

- Log in to the AWS Management Console and navigate to the DynamoDB service.
- 2. In the left-hand menu, click on **Tables**.
- 3. Click the Create table button.
- 4. Under Table settings:
 - Table name: Enter bookstore.
 - Partition key:
 - Name: id
 - **Type:** Select Number from the dropdown.
- 5. Keep all other settings as **default** (for example, on-demand capacity, encryption, etc.).
- 6. Click Create table.
- 7. Wait until the table's status becomes **Active**.

Step 2: Create a Lambda Function

2.1 Create a Lambda Function Using a Blueprint

- 1. In the AWS Management Console, navigate to **Lambda**.
- 2. Click on Create function.
- 3. Under Create function, select Use a blueprint.
- In the blueprint search field, type and select "Create a mobile backend that interacts with a DDB table".
 - This blueprint is preconfigured to interact with DynamoDB.
- 5. Under Basic information:
 - Function name: Enter testing1.
 - Runtime: (The blueprint will auto-select an appropriate runtime, for example, Python or Node.js.)
- Under Permissions, choose Create a new role with basic Lambda permissions.
- 7. Click Create Function.

2.2 Configure and Test the Lambda Function

1. After creation, click on the **Code** tab in your Lambda function's details page.

2. Configure a Test Event:

- Click on Test.
- o In the pop-up, click on **Configure test event**.
- Select Create new event.
- o **Event name:** Enter testing001.

In the **Event JSON** section, paste sample JSON data for a book record. For example:

```
"operation": "create",

"data": {

    "id": 101,

    "title": "Serverless Essentials",

    "author": "AWS User",

    "price": 29.99
}
```

- Click Create.
- 3. Attach a Policy for DynamoDB Access:

- Switch to the **Configuration** tab of the Lambda function.
- Under Permissions, click on the Role name link. This will open the
 IAM role associated with your function.
- In the IAM role page, click Attach policies.
- Search for AmazonDynamoDBFullAccess.
- Check the box next to it and click Attach policy.
 - This gives the Lambda function full access to DynamoDB.
- 4. Return to the Lambda function's **Code** tab.
- 5. Click **Test** (using the testing001 event) to invoke the function.
- Open the **DynamoDB Console**, select your **bookstore** table, and click on **Explore items**.
 - Verify that the test event data has been inserted into the table.
 - (The blueprint's code typically performs an operation like inserting the JSON data into the table.)

Step 3: Test with API Gateway

3.1 Create a REST API in API Gateway

- 1. In the AWS Management Console, navigate to **API Gateway**.
- 2. Click on **REST API**.

- 3. Choose New API.
- 4. API Name: Enter bookstore.
- 5. Click Create API.

3.2 Create a Resource in the API

- 1. In the newly created API, click on **Actions** and then **Create Resource**.
- 2. For **Resource Name**, enter bookid.
- 3. For **Resource Path**, enter / (if you want the root resource, you can simply name it bookid for identification).
 - Alternatively, you can create a subresource / bookid.
- 4. Click Create Resource.

3.3 Create a PUT Method for the /bookid Resource

- With the /bookid resource selected, click on Actions and then Create
 Method.
- 2. Choose the **PUT** method from the dropdown.
- 3. Click the checkmark to confirm.
- 4. In the **PUT setup**:
 - Integration type: Select Lambda Function.
 - Use Lambda Proxy integration: (Typically enabled; if not, you may need to configure mapping templates.)

- Lambda Region: Ensure it is set to the same region as your Lambda function.
- Lambda Function: Enter or select testing1 (the Lambda function you created earlier).
- Click Save, and then acknowledge any prompts regarding Lambda permissions.
- After the method is created, click on **Test** within the PUT method configuration.

In the test window, paste any random JSON data for a book record similar to:

```
"operation": "create",

"data": {

   "id": 102,

   "title": "AWS Serverless Cookbook",

   "author": "DevOps Guru",

   "price": 39.99
}
```

}

- 7. Click **Test** and verify that the test invocation is successful.
- 8. Switch to the **DynamoDB Console**, select your bookstore table, and confirm that the new data is visible in **Explore items**.

3.4 Deploy the API

- 1. In API Gateway, click on **Actions → Deploy API**.
- 2. In the **Deployment stage** dropdown, choose [New Stage].
- 3. Enter a **Stage Name** (e.g., prod).
- 4. Click **Deploy**.
- Note the **Invoke URL** provided after deployment. This URL is now accessible publicly.

Step 4: Final Verification

4.1 Using MySQL Workbench/Query Tools (Optional)

 Although this step is more focused on API testing, you can also manually verify data by querying the DynamoDB table.

4.2 Verify the End-to-End Flow

1. Invoke the API:

 Use a tool like **Postman** or **curl** to send a PUT request to your API's /bookid resource.

```
Example using curl:
```

```
curl -X PUT \
     -H "Content-Type: application/json" \
     -d '{
           "operation": "create",
           "data": {
             "id": 103,
             "title": "Serverless in Action",
             "author": "AWS Guru",
             "price": 49.99
           }
         }' \
```

https://<your-api-id>.execute-api.<region>.amazonaws.com/pro
d/bookid

2. Check DynamoDB:

- Go to the **DynamoDB Console** and open the bookstore table.
- o Confirm that the new data (id 103) has been inserted.

3. Test via Lambda Console:

 You may also run additional test events from the Lambda console to confirm the function's behavior.



What to Expect After Implementing This Project

Task?

By completing this hands-on AWS Serverless Architecture project, you will have successfully built a fully functional, scalable, and cost-efficient backend solution. Here's what you can expect after implementation:

* Key Achievements:

✓ A Fully Serverless Backend

- You will have built a completely serverless application using AWS services—Lambda, API Gateway, and DynamoDB—without provisioning or managing any servers.
- This architecture auto-scales based on traffic and runs only when needed, helping optimize cost and performance.
- V DynamoDB as a Highly Available NoSQL Database

- Your DynamoDB table (bookstore) will dynamically store, retrieve, and manage book records with low latency and high availability.
- It offers on-demand scaling and eliminates the need for database administration.

AWS Lambda for Event-Driven Computing

- Your Lambda function (testing1) will process API requests, execute backend logic, and interact with DynamoDB—all without running continuously.
- This results in **cost efficiency**, as you only pay for execution time.

✓ API Gateway for RESTful Endpoints

- Your API Gateway instance (bookstore API) will act as the entry point for external applications, allowing them to interact with your DynamoDB table using simple HTTP methods like PUT, GET, DELETE, and POST.
- API Gateway ensures secure and managed API access with built-in logging and monitoring.

▼ End-to-End Integration & Testing

- You will have tested the entire pipeline using Postman, cURL, and AWS
 Console to validate database interactions.
- You will have successfully inserted and retrieved book records in DynamoDB via API requests.

Why This Project is Valuable?

Real-World AWS Experience

- This is a practical, industry-relevant AWS project that mirrors real-world scenarios used by companies for scalable, event-driven applications.
- Ideal for DevOps, cloud computing, and backend development roles.

S Cost Optimization

- Unlike traditional server-based architectures, this setup ensures that you only pay for actual usage, making it highly cost-efficient.
- No need to manage infrastructure, reducing operational overhead.

🔒 Highly Scalable & Secure

- AWS auto-scales Lambda and DynamoDB based on demand.
- API Gateway provides authentication, rate limiting, and monitoring for enhanced security.

% What's Next?

Now that you have successfully implemented the **AWS Serverless Architecture**, you can further extend and optimize the project:

Enhancements & Advanced Features:

- Add More API Methods Implement GET, DELETE, and UPDATE
 endpoints for a complete CRUD (Create, Read, Update, Delete) API.
- Implement Authentication Secure your API with AWS Cognito for user authentication.
- Optimize Performance Use DynamoDB Streams to trigger Lambda functions on data changes.
- Monitor with AWS CloudWatch Set up logging, metrics, and alerts for better observability.
- Deploy with Infrastructure as Code Automate deployments using
 Terraform or AWS CloudFormation.

Summary

DynamoDB Table Creation:

 A table named bookstore is created with a partition key id (Number).

• Lambda Function Setup:

- A Lambda function (testing1) is created using a blueprint that interacts with DynamoDB.
- A test event is configured, and the function is granted full access to
 DynamoDB via an attached policy.

• API Gateway Configuration:

- o A new REST API (bookstore) is created.
- A resource /bookid is set up with a PUT method integrated with the Lambda function.
- The API is deployed to a stage (e.g., prod).

• End-to-End Testing:

- Test data is sent via API Gateway, which triggers the Lambda function to insert data into DynamoDB.
- The inserted data is verified in the DynamoDB console.

Wrapping Up: Mastering Serverless

Architecture with AWS!

Congratulations on reaching the end of this comprehensive guide on AWS

Serverless Architecture! We've covered everything from theory to real-world implementation, ensuring that you now have a solid understanding of AWS

Lambda, API Gateway, and DynamoDB in action.

- You've learned how serverless architecture eliminates infrastructure management and enables scalability, cost-efficiency, and event-driven workflows.
- You've built a fully functional API that seamlessly interacts with DynamoDB via AWS Lambda.

 You've explored best practices for designing, deploying, and securing serverless applications.



What's Next?

AWS offers an entire **ecosystem of serverless services**, and we've only scratched the surface. As we move forward, I'll be covering **more advanced AWS services** with real-time implementation projects.

Daily AWS Practical Tasks Incoming!

I'm committed to sharing **AWS real-world projects** daily, covering different services with hands-on implementation. **Follow me to stay updated and keep learning!**